



UNIVERSITÀ DI PISA

A.Y. 2019/20

« Intelligent systems for pattern recognition »

Master Degree in Computer Science

Artificial Intelligence Curriculum

*Midterm 3*  
*Assignment 3*

---

*Gated Recurrent  
Neural Networks  
(LSTM)*

*Diletta Goglia*

# Data loading and preprocessing

## Model 1

- Multivariate timeseries (multiple features) with multivariate input data
- 18 sensors in the house, in 9 different rooms, both for temperature and humidity.
- 18 features in input, 1 feature for output ("Appliances" column)
- Measurements from January to April for training
- Month of May for test
- `sklearn.preprocessing` library to scale data

## Model 2

- Load only the "appliances" column from dataset
- Different train/test splitting
- Different input/output
- Input shape: passing only one step in the past
- Reshaping the output is not needed (because `return_sequences` in LSTM layer is set to `False`).

```
87
88 ''' Model 1) consider temperature and humidity data as input '''
89 # multivariate input data
90 # Load CSV with Pandas
91 dataset = pd.read_csv('energydata_complete.csv', header = 0, sep=',', quote=
92
93 #datetime format
94 dataset.index = pd.to_datetime(dataset['date'], format='%Y-%m-%d %H:%M:%S')
95 dataset = dataset.set_index('date')
100
110 ''' DATASET SPLITTING '''
111
112 train = dataset["2016-01-11":"2016-04-30"] # jan - apr
113 test = dataset["2016-05-11":"2016-05-27"] # may
114
115 train.index = pd.to_datetime(train.index, format='%Y-%m-%d %H:%M:%S')
116 test.index = pd.to_datetime(test.index, format='%Y-%m-%d %H:%M:%S')
117
118 training_values = train.values
119 test_values = test.values
120
135 ''' INPUT - OUTPUT '''
136 # split into input and outputs
137 train_X, train_y = training_values[:, 1:], training_values[:, :1]
138 # temperature e umidità (9+9 tot. 18 feature) sono input
139 # mentre la prima colonna (appliances) è l'output
140 test_X, test_y = test_values[:, 1:], test_values[:, :1]
141
142
143 ''' PREPROCESSING '''
144 # normalize features
145 scaler = MinMaxScaler(feature_range=(0, 1))
146 scaler.fit_transform(train_X)
147 train_X = scaler.transform(train_X)
148 scaler.fit_transform(train_y)
149 train_y = scaler.transform(train_y)
150 scaler.fit_transform(test_X)
151 test_X = scaler.transform(test_X)
152 scaler.fit_transform(test_y)
153 test_y = scaler.transform(test_y)
154
155 # reshape input to be 3D [samples, timesteps, features]
156 train_X = train_X.reshape(train_X.shape[0], 1, train_X.shape[1])
157 test_X = test_X.reshape(test_X.shape[0], 1, test_X.shape[1])
158
```

## Train / Test

	Appliances	T1	...	T9	RH_9	
date						
2016-01-11 17:00:00	60.0	19.890000	...	17.033333	45.5300	Train
2016-01-11 17:10:00	60.0	19.890000	...	17.066667	45.5600	
2016-01-11 17:20:00	50.0	19.890000	...	17.000000	45.5000	
2016-01-11 17:30:00	50.0	19.890000	...	17.000000	45.4000	
2016-01-11 17:40:00	60.0	19.890000	...	17.000000	45.4000	
...						
2016-05-27 17:20:00	100.0	25.566667	...	23.200000	46.7900	Test
2016-05-27 17:30:00	90.0	25.500000	...	23.200000	46.7900	
2016-05-27 17:40:00	270.0	25.500000	...	23.200000	46.7900	
2016-05-27 17:50:00	420.0	25.500000	...	23.200000	46.8175	
2016-05-27 18:00:00	430.0	25.500000	...	23.200000	46.8450	

[19735 rows x 19 columns]

## Input / Output

	Output	Input				Features
	Appliances	T1	...	T9	RH_9	
date						
2016-01-11 17:00:00	60.0	19.890000	...	17.033333	45.5300	Train
2016-01-11 17:10:00	60.0	19.890000	...	17.066667	45.5600	
2016-01-11 17:20:00	50.0	19.890000	...	17.000000	45.5000	
2016-01-11 17:30:00	50.0	19.890000	...	17.000000	45.4000	
2016-01-11 17:40:00	60.0	19.890000	...	17.000000	45.4000	
...						
2016-05-27 17:20:00	100.0	25.566667	...	23.200000	46.7900	Test
2016-05-27 17:30:00	90.0	25.500000	...	23.200000	46.7900	
2016-05-27 17:40:00	270.0	25.500000	...	23.200000	46.7900	
2016-05-27 17:50:00	420.0	25.500000	...	23.200000	46.8175	
2016-05-27 18:00:00	430.0	25.500000	...	23.200000	46.8450	

[19735 rows x 19 columns]

## 3D tensors reshape

Train\_input (temperature and humidity)  
(15882, 18)  
Train\_output (energy)  
(15882, 1)  
Test\_input (temperature and humidity)  
(2413, 18)  
Test\_output (energy)  
(2413, 1)

After reshaping:  
(15882, 1, 18) (15882, 1) (2413, 1, 18) (2413, 1)

# Code snippets

- Keras Sequential API
- Build LSTM
- Add layers
- Model fit
- Check loss
- Predict

## Model 1

```
163 ''' HYPERPARAM '''
164 learning_rate=1e-4
165 batch_size=70
166 epochs=20
167
168 ''' BUILDING THE FIRST MODEL '''
169 # using Sequential API
170 model = tf.keras.Sequential() # instaciate a model using Sequential class
171 # --> will contruct a pipeline of layers
172 # building add one layer at time
173 model.add(layers.LSTM(18, activation='tanh', return_sequences=True,
174                       input_dim=(train_X.shape[2]))) 18 input features
175 # set the return_sequences to True, the output shape becomes a 3D array
176 model.add(layers.Dropout(0.5))
177 # Dropout regularize the model by randomly tuting off some neurons
178 # --> prevent overfitting
179 model.add(layers.Dense(1))
180 # then you don't need to specify the input shape again because
181 # it is automatically inferred by sequential layer
182
183 model.compile(optimizer='adam',
184               loss='mae',
185               metrics=['accuracy'])
186
187 print('\n', model.summary())
188
189 ''' MODEL FIT '''
190 # fit model
191 history = model.fit(train_X, train_y, validation_split=0.2, epochs=epochs,
192                    batch_size=batch_size, verbose=1, shuffle=False)
193
194 # Plot Model Loss
195 # list all data in history
196 print(history.history.keys())
197 # summarize history for loss
198 plt.figure(3)
199 plt.plot(history.history['loss'])
200 plt.plot(history.history['val_loss']) #RAISE ERROR
201 plt.title('model loss')
202 plt.ylabel('loss')
203 plt.xlabel('epoch')
204 plt.legend(['train', 'test'], loc='upper left')
205 plt.show()
206
207 ''' MODEL PREDICTION '''
208 # make a prediction
209 test_predict0 = model.predict(test_X)
210 #reshape (because return_sequences was set to True)
211 test_predict = test_predict0.reshape((test_predict0.shape[0],
212                                     test_predict0.shape[2]))
213 # invert predictions
214 test_predict = scaler.inverse_transform(test_predict)
```

## Model 2

pretty much the same, but...

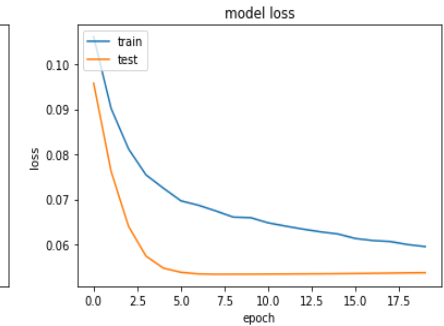
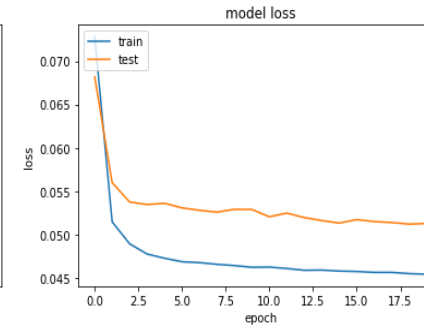
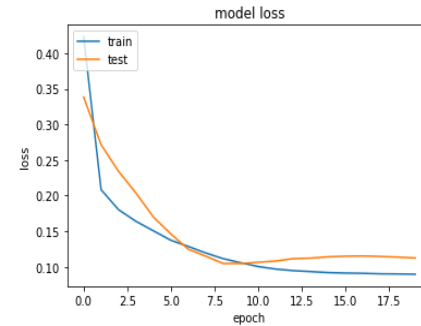
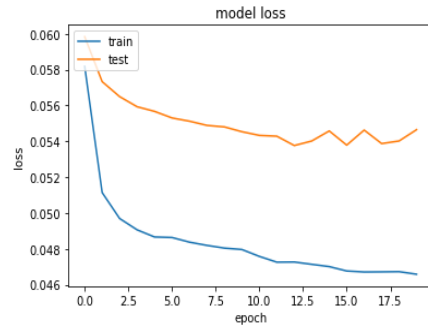
```
115 ''' PREPROCESSING'''
116
117 scaler = MinMaxScaler(feature_range=(0, 1))
118 dataset = scaler.fit_transform(dataset)
119 train_size = int(len(dataset) * 0.80)
120 test_size = len(dataset) - train_size
121 train1, test1 = dataset[0:train_size,:], dataset[train_size:len(dataset),:]
122
123 def create_dataset(dataset, look_back=1):
124     X, Y = [], []
125     for i in range(len(dataset)-look_back-1):
126         a = dataset[i:(i+look_back), 0]
127         X.append(a)
128         Y.append(dataset[i + look_back, 0])
129     return np.array(X), np.array(Y)
130
131 ''' DATASET SPLITTING '''
132
133 look_back = 1
134 X_train, Y_train = create_dataset(train1, look_back)
135 X_test, Y_test = create_dataset(test1, look_back)
136
137 # reshape input to be [samples, time steps, features]
138 X_train = np.reshape(X_train, (X_train.shape[0], 1, X_train.shape[1]))
139 X_test = np.reshape(X_test, (X_test.shape[0], 1, X_test.shape[1]))
140
141 ''' BUILDING THE SECOND MODEL '''
142
143 # using Sequential API
144 model = tf.keras.Sequential() # instaciate a model using Sequential class
145 # --> will contruct a pipeline of layers
146 model.add(layers.LSTM(10, input_shape=(X_train.shape[1], X_train.shape[2]))) 1 timestep, 1 input feature
147 model.add(layers.Dropout(0.2))
148 model.add(layers.Dense(1))
149
150 #model.compile(loss='mean_squared_error', optimizer='adam')
151 model.compile(optimizer='adam', # try also adam
152               loss='mae')
153
154 ''' MODEL FIT '''
155 # fit model
156 history = model.fit(X_train, Y_train, epochs=epochs, batch_size=batch_size, v
157                    verbose=1, shuffle=False)
158 model.summary()
```

# Model 1:

Predict the current energy expenditure given as input information the temperature ( $T_i$ ) and humidity ( $RH_i$ ) information from all the  $i$  sensors in the house.

## Tuning hyperparameters

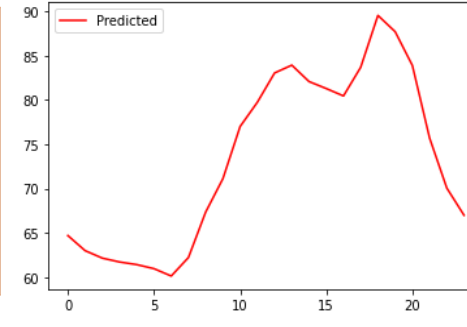
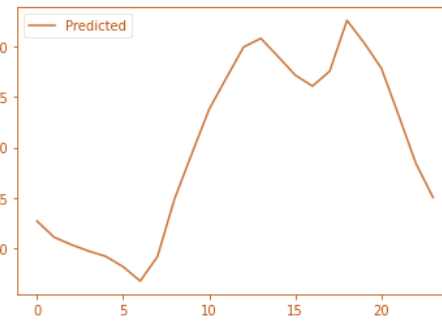
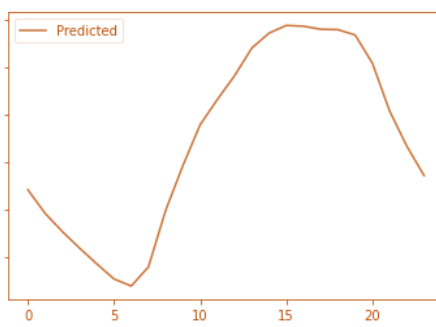
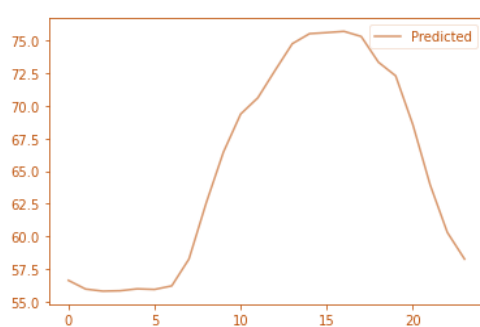
- Dropout
- Number of hidden layers
- Number of neurons for LSTM layer
- Optimizers
- Learning rate
- Loss functions
- Activation functions
- Batch size
- Range of feature scaling



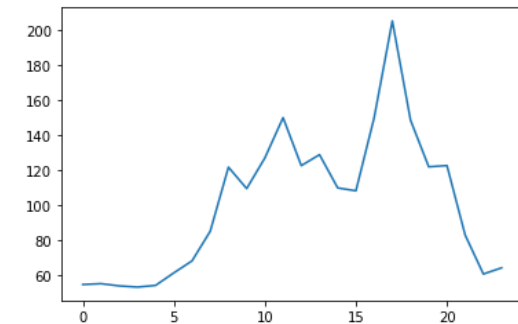
## Results

\* Plotted grouping by mean value of energy consumption

### Predictions



### Real data

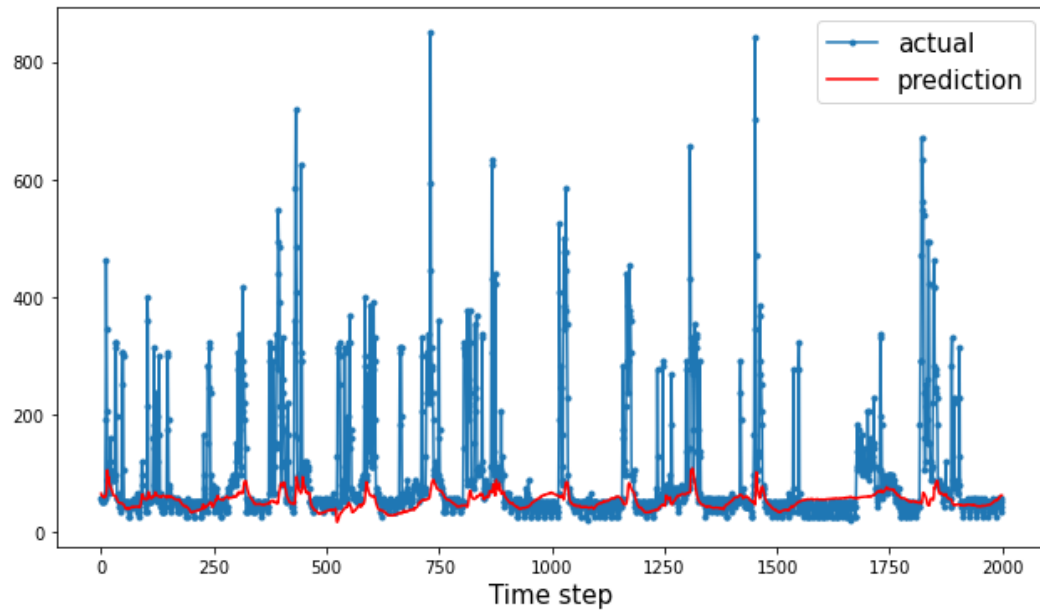


The model is able to capture the overall trend.

# *Model 1:*

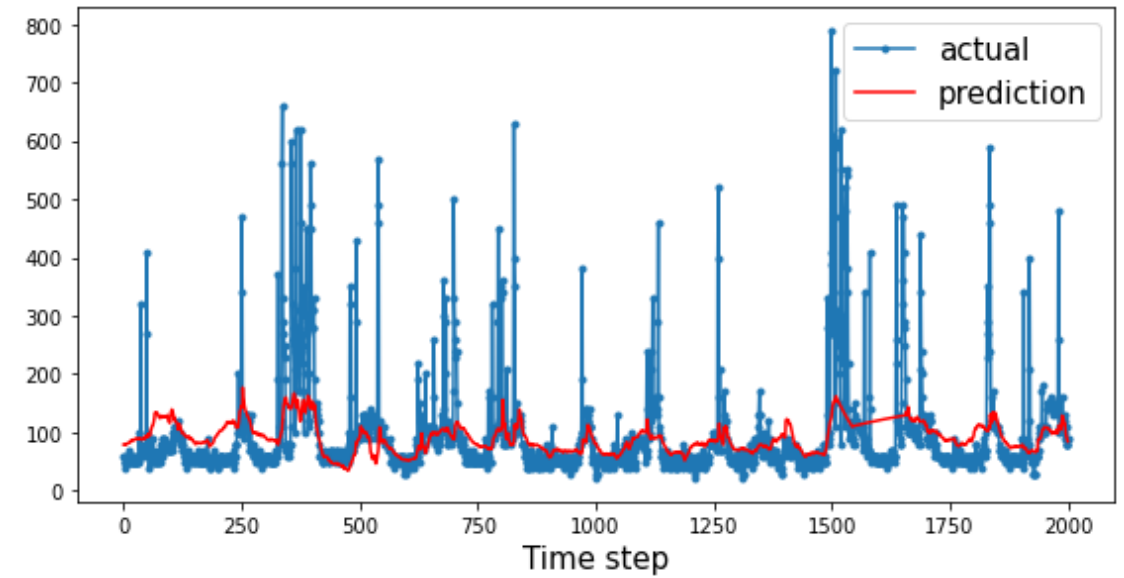
## *Predicted values VS Real Data*

*Validation*



Predicted values are improved by tuning the model parameters and, at the end, they are almost exactly like ground truth data.

*Test*

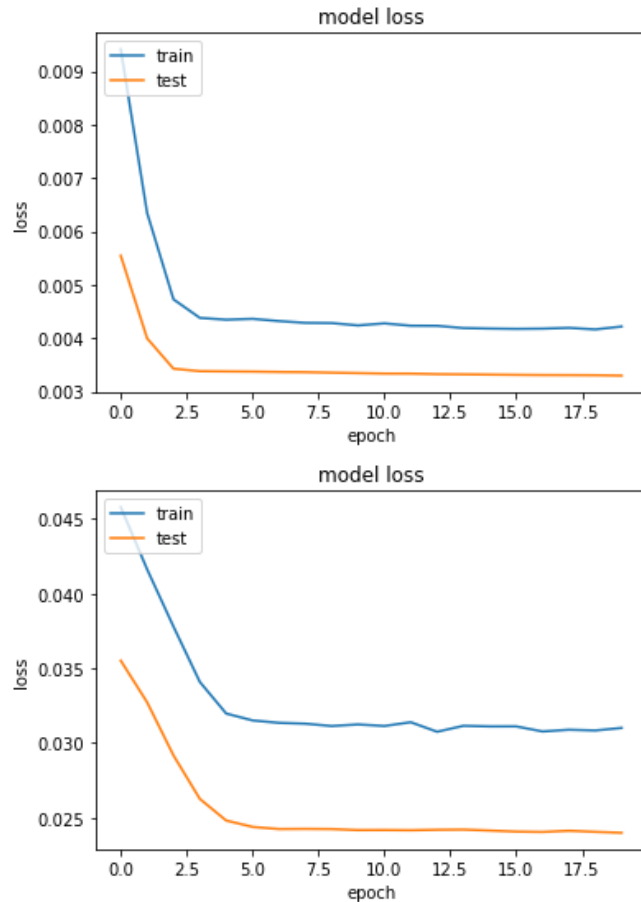


Not perfect... but really good prediction!

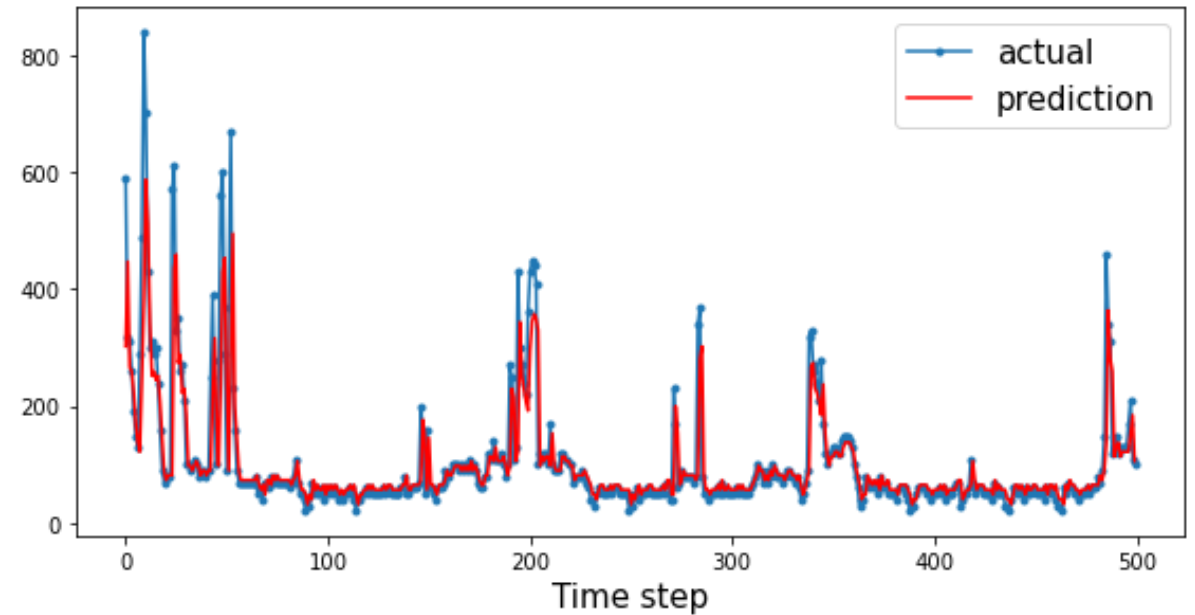
# Model 2:

Setup a one step-ahead predictor for energy expenditure, i.e. given the current energy consumption, predict its next value.  
Predict the value at the current time step by using the history ( $n$  time steps from it, in this case, with  $n=1$ )

## Tuning hyperparameters



## Results



# Final considerations

## Pros and cons of Teacher Forcing (used in the first model)

### Pros:

- Training with *Teacher Forcing* converges faster. If we do not use *Teacher Forcing*, the hidden states of the model will be updated by a sequence of wrong predictions, errors will accumulate, and it is difficult for the model to learn from that.

### Cons:

- During inference, since there is usually no ground truth available, the model will need to feed its own previous prediction back to itself for the next prediction. Therefore there is a discrepancy between training and inference, and this might lead to poor model performance and instability. This is known as *Exposure Bias* in literature.

“... [teacher forcing] can result in problems in generation as small prediction error compound in the conditioning context. This can lead to poor prediction performance as the RNN’s conditioning context (the sequence of previously generated samples) diverge from sequences seen during training.”

“The disadvantage of strict teacher forcing arises [...] the fed-back inputs that the network sees during training could be quite different from the kind of inputs that it will see at test time.”

## Why predictions in the second model are better?

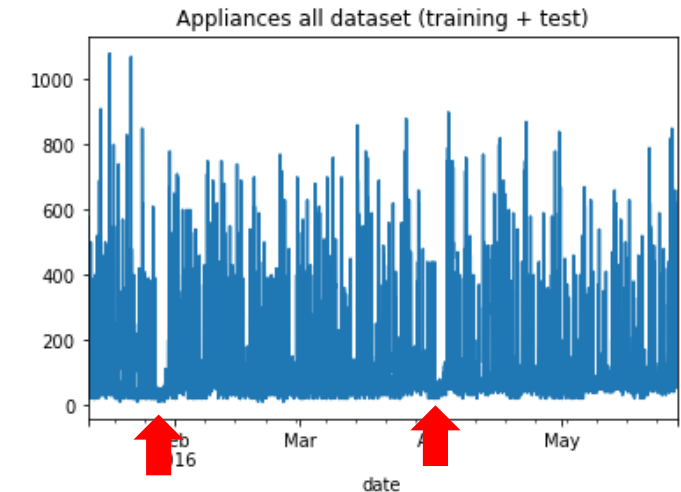
- Second model has much less parameters to train
- Lag in appliances dataset
- Not all the temperatures and humidity measurement must be related to appliances consumption!

## Time series correlations

	Appliances	T1	RH_1	...	RH_8	T9	RH_9
Appliances	1.000000	0.055447	0.086031	...	-0.094039	0.010010	-0.051462
T1	0.055447	1.000000	0.164006	...	-0.006441	0.844777	0.071756
RH_1	0.086031	0.164006	1.000000	...	0.736196	0.115263	0.764001
T2	0.120073	0.836834	0.269839	...	0.068534	0.675535	0.157346
RH_2	-0.060465	-0.002509	0.797535	...	0.679777	0.054544	0.676467
T3	0.085060	0.892402	0.253230	...	0.044427	0.901324	0.134602
RH_3	0.036292	-0.028550	0.844677	...	0.828822	-0.195270	0.833538
T4	0.040281	0.877001	0.106180	...	-0.095192	0.889439	-0.025549
RH_4	0.016965	0.097861	0.880359	...	0.847259	-0.044518	0.856591
T5	0.019760	0.885247	0.205797	...	0.016388	0.911055	0.072308
RH_5	0.006955	-0.014782	0.303258	...	0.359840	-0.138509	0.272197
T6	0.117638	0.654769	0.316141	...	0.073721	0.667177	0.184424
RH_6	-0.083178	-0.615045	0.245126	...	0.489580	-0.738940	0.391943
T7	0.025801	0.838705	0.021397	...	-0.209961	0.944776	-0.077690
RH_7	-0.055642	0.135182	0.801122	...	0.883984	0.028055	0.858686
T8	0.039572	0.825413	-0.030053	...	-0.209532	0.869338	-0.156820
RH_8	-0.094039	-0.006441	0.736196	...	1.000000	-0.113014	0.855812
T9	0.010010	0.844777	0.115263	...	-0.113014	1.000000	-0.008683
RH_9	-0.051462	0.071756	0.764001	...	0.855812	-0.008683	1.000000

[19 rows x 19 columns]

\* the correlation does not depend on a cause-effect relationship but on the tendency of one variable to change according to another



Almost NOT related!

*Thanks for your  
attention*



UNIVERSITÀ DI PISA

A.Y. 2019/20

« Intelligent systems for pattern recognition »

Master Degree in Computer Science

Artificial Intelligence Curriculum

---